





Graph Contractions in Redstar

Oguz Selvitopi, Aydın Buluç (LBNL)
Emin Ozturk, P. (Saday) Sadayappan (Utah)
Jie Chen, Robert Edwards (JLab)


Computational phases & graphs



**Graph
isomorphism**



**Contraction
path trees**




Contraction graphs

**Scheduling
Partitioning**

Contraction queue &
batched contractions



**Contraction
DAG**



**Graph
reordering**

Scheduling


▷ Scheduling

- Optimize memory utilization on **single** GPU
- **Goal**: Increase tensor reuse
 - Reduce evictions, data transfer between host and device
- **How**: Reorder contractions (contraction DAG)

▷ Three heuristics

- Sibling-based
- Node-based
- Tree-based

Memory model and scheduling order



n contractions c_1, c_2, \dots, c_n

Memory in use after contraction c_i : M_i

Peak memory: $\max_i M_i$

Order	Contents	Size
e	{b, e}	2
g	{a, b, e}	3
h	{a, b}	2
f	{}	0

Order #1

Peak memory: 3


Order	Contents	Size
f	{a, b}	2
e	{a, e}	2
g	{a}	1
h	{}	0

Order #2

Peak memory: 2


Sibling-based scheduler

- ▶ Exploit the specific property that **each contraction is binary**
 - Each node \rightarrow two children
- ▶ Idea: After completing a contraction, process its **sibling** contractions
- ▶ Motivation: Enable contractions higher in the contraction DAG
 - Higher priority for the contractions higher in the DAG
 - Achieve a DFS scheduling of contractions to reduce memory footprint



Node-based scheduler

- ▶ More **general** scheduling choices
 - Do not depend on specific structure of the DAG
- ▶ **Idea**: Choose the contraction that causes **least amount of increase in memory**
 - Most recent state of the memory
- ▶ **Motivation**: Scheduling decisions based on **an objective to reduce memory footprint**
 - For each node **u**, maintain
 - **$u.\delta$** : change in utilized memory
 - Global view of the DAG
 - Choose a contraction among all that can be scheduled




GRAY Completed contractions / ready tensors

GREEN Can be scheduled

RED Depend on tensors not yet available

Tree-based scheduler

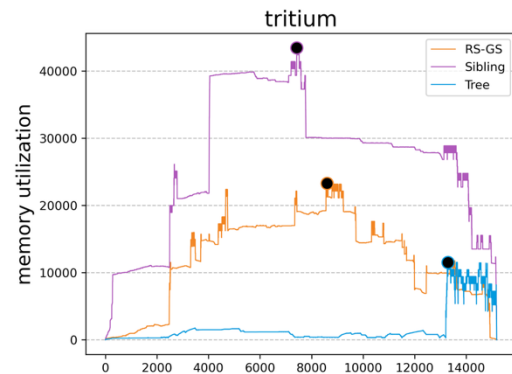
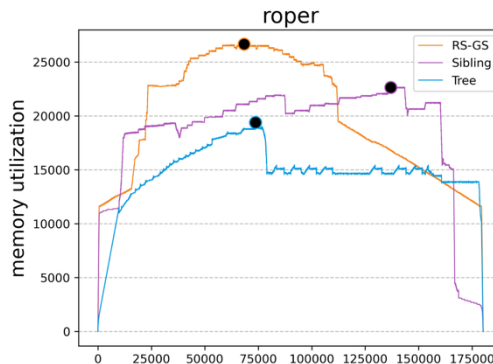



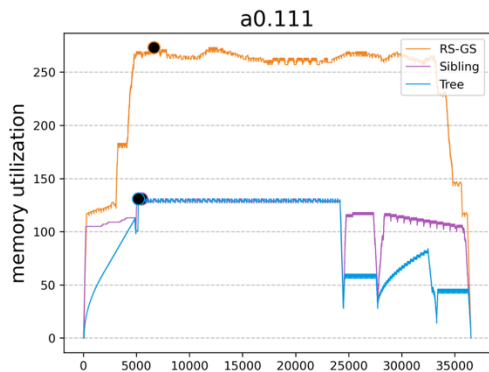
- ▶ Similar to node-based scheduler but:
 - Schedule **a subset of nodes** instead of a single node
 - Subset of nodes
 - \equiv nodes in contraction trees
- ▶ **Motivation**
 - Nodes in contraction trees are connected
 - Inherent locality
- ▶ For each tree T_i
 - **gain(T_i)**: Change in memory if contractions in T_i were performed
- ▶ Select T_i with **highest gain among all trees**
 - \equiv smallest increase in memory



Peak memory

a0.111: MxM (18K vertices, 36K edges)
a0.d3: MxM (3.8K vertices, 7.2K edges)
f0: MxMxM (30K vertices, 59K vertices)

roper: BxM (90K vertices, 180K edges)
deuteron: BxB (156K vertices, 312K edges)
tritium: BxBxB (7.5K vertices, 15K vertices)



Data movement

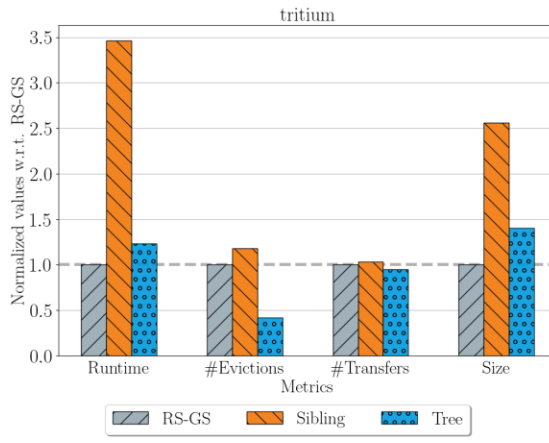

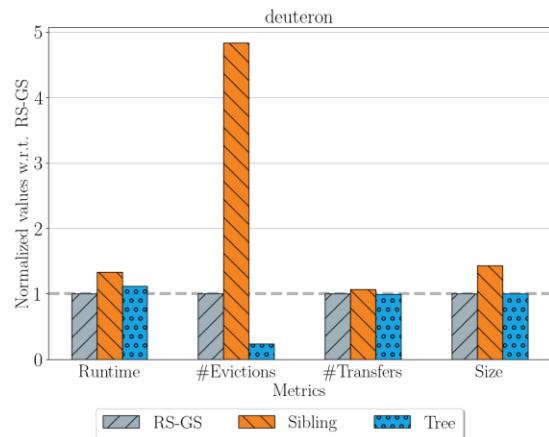
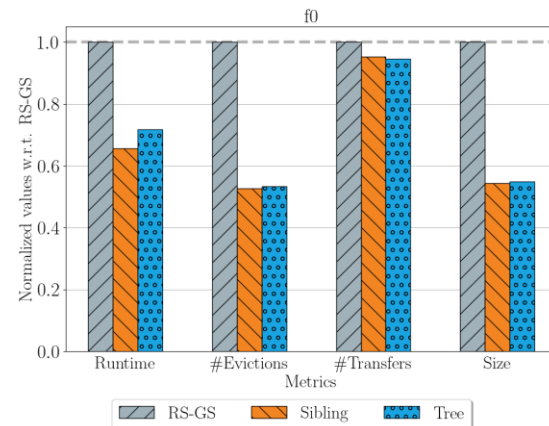


Corr. Func.	RS-GS	Sibling	Tree
a0-111	2.00	1.12	1.11
a0-d3	1.12	0.84	0.88
f0	1.01	0.55	0.55
roper	7.35	4.91	5.47
deuteron	0.21	0.29	0.20
tritium	0.53	1.13	0.74

Size in TBs

Corr. Func.	RS-GS	Sibling	Tree
a0-111	3.5	29.5	189.1
a0-d3	0.5	3.8	19.2
f0	5.9	50.0	295.2
roper	26.0	234.3	3095.9
deuteron	38.4	451.0	17005.5
tritium	1.5	25.0	212.5


Time in msec

Redstar runtime



Partitioning

- ▶ Distribute contractions among GPUs
 - **Goal:** Reduce data transfers
 - **h2d:** **slow** (PCIe)
 - **d2d:** **fast** (NVLink)
 - Balance GPU loads
- ▶ Contraction DAG
 - **Leaves** → h2d
 - **Non-leaves** → d2d



Partitioning model

▷ Desired


- **#1** Roughly equal amount of contractions from each level
- **#2** Reduction of h2d is more important than reduction of d2d

▷ Model

- Each contraction/tensor \rightarrow vertex
- Edge between vertices \rightarrow dependency of a contraction on a tensor


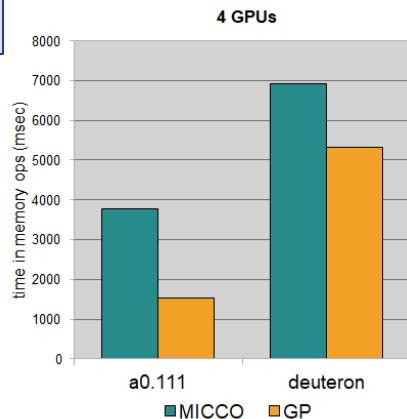
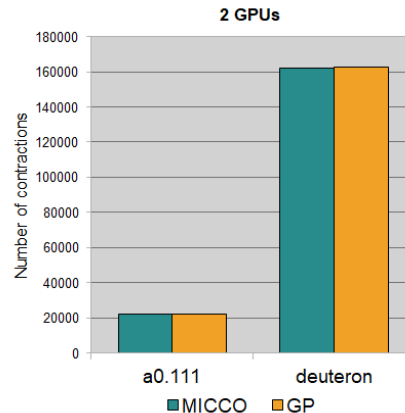

▷ Vertex weights

- **$c = \text{number of levels} - 1$** \equiv number of constraints
- vertex at level i
 - $w_i(v) = \text{contraction cost}$, $w_{1 \leq j \neq i \leq c}(v) = 0$
- leaves \rightarrow no computation, no weight



Boundary replication

No communication among devices
= zero d2d operations (**GP**)
Still communication between host and device



Replicate contractions on the boundary
= extra work on GPUs (**GP**)